

How to find a cozy room?

A data mining project in rental market



Warm-up

- Have you ever rented a room?
- What factors will you consider?
- Will you take the room shown here?



4500 元/月 (季付价)

近地铁 随时看房

整租
 1室0厅1卫
 33m²
 朝南



Aim

- In this project, we want to study the rental market in Shanghai in a quantitative way by investigating the relationship between unit price and some other related variables.



Plan

- Crawl the data from <http://sh.lianjia.com/zufang>
- Clean the data and extract factors, unit price and more
- Visualize the factors
- Explore the relationship between unit price and variables

首页 二手房 新房 租房 海外 小区 经纪人 指南 工具 发布房源

链家 首页 整租 合租 下载APP

请输入区域、商圈或小区名开始找房

链家网上海站 > 上海租房

按区域 ▲ 按地铁线 ▼

不限 静安 徐汇 黄浦 长宁 普陀 浦东 宝山 闸北 虹口 杨浦 阳行 金山 嘉定 崇明 奉贤 松江 青浦

方式 不限 整租 合租

租金 ≤1000元 1000-1500元 1500-2000元 2000-3000元 3000-5000元 5000-8000元 ≥8000元 - 元 确定

户型 一居 两居 三居 四居+

朝向 东 西 南 北 南北

更多 ▼

已为您找到 23968 套 上海租房 清空条件

综合排序 最新上架 价格 面积

整租 · 武定西路1371弄 1室0厅 4500元 4500 元/月

长宁-镇宁路 / 33m² / 南 / 1室0厅1卫


链家 ◎ 1个月前发布

近地铁 随时看房

整租 · 花苑茶花园 2居室 5000

徐汇-康健 / 59m² / 南 / 2室1厅1卫

下载
扫描上方二维码
随时随地查看
了解更多



Tools

- Chrome: learn to press F12
- Python with packages
- Web crawler: urllib, requests, re, beatifulsoup
- Data analysis: NumPy, pandas, scikit-learn (sklearn)
- Data visualization: pyecharts, matplotlib



Web crawler

- A Web crawler starts with **a list of URLs to visit**, called the *seeds*. As the crawler visits these URLs, it identifies all the hyperlinks in the page and **adds them to the list of URLs to visit**, called the crawl frontier. URLs from the frontier are recursively visited according to a set of policies. If the crawler is performing archiving of websites it saves the information as it goes. The archives are usually stored as 'snapshots'.

Web crawler


- Situations to face
 - The archive is known as the *repository* and is designed to store and manage the collection of web pages. The repository only stores HTML pages and these pages are stored as distinct files.
 - The large volume implies the crawler can only download a limited number of the Web pages within a given time.
 - The number of possible URLs crawled being generated by server-side software has also made it difficult for web crawlers to avoid retrieving duplicate content.

Web crawler

- Crawling policy
 - a *selection policy* which states the pages to download,
 - a *re-visit policy* which states when to check for changes to the pages,
 - a *politeness policy* that states how to avoid overloading Web sites.
 - a *parallelization policy* that states how to coordinate distributed web crawlers.



Web crawler



Web crawler: html

- **Hypertext Markup Language (HTML)**
 - is the standard [markup language](#) for creating web pages and web applications. With [Cascading Style Sheets](#) (CSS) and [JavaScript](#), it forms a triad of cornerstone technologies for the World Wide Web.
 - Web browsers receive HTML documents from a [web server](#) or from local storage and [render](#) the documents into multimedia web pages. HTML describes the structure of a web page [semantically](#) and originally included cues for the appearance of the document.

Web crawler: html

- HTML elements
 - building blocks of HTML pages. Images and objects such as interactive forms may be embedded into the page.
 - text can have headings, paragraphs, lists, links, quotes and other items.
 - HTML elements are denoted by *tags*, written using angle brackets. Tags such as `` and `<input />` directly introduce content into the page.

Web crawler: web pages

已为您找到 23968 套 上海租房

Press F12 in Chrome

清空条件

整租 · 武定西路1371弄 1室0厅 4500元

长宁-镇宁路 / 33m² / 南 / 1室0厅1卫
链家 ⚡ 1个月前发布

4500 元/月

近地铁 随时看房

整租 · 花苑茶花园 2居室 5000

徐汇-康健 / 59m² / 南 / 2室1厅1卫
门店优选 ⚡ 5个月前发布

5000 元/月

随时看房

整租 · 近8号线杨思站 简装两房 诚意出租

浦东-三林 / 61m² / 南 / 2室1厅1卫
链家 ⚡ 1个月前发布

4800 元/月

近地铁 随时看房

整租 · 房东诚意出租 看房有钥匙 价格还能谈

徐汇-斜土路 / 57m² / 南北 / 2室1厅1卫
链家 ⚡ 1个月前发布

6500 元/月

随时看房

Elements Console Sources Network Performance Memory Application Security Audits

Styles Computed Event Listeners

Filter element.style {

.content index.css? v=20.411190841403:1
.content .content_list {
height: 182px;
margin-top: 40px;
position: relative;
}

a, abbr, common.css? v=2.411190841403:1
acronym, address, applet, article, aside,
audio, b, big, blockquote, body, canvas,
caption, center, cite, code, dd, del,
details, dfn, div, dl, dt, em, embed,
fieldset, figcaption, figure, footer, form,
h1, h2, h3, h4, h5, h6, header, hgroup, html,
i, iframe, img, ins, kbd, label, legend, li,
mark, menu, nav, object, ol, output, p, pre,
q, ruby, s, samp, section, small, span,
strike, strong, sub, summary, sup, table,
tbody, td, tfoot, th, thead, time, tr, tt, u,
ul, var, video {
margin: 0;
padding: 0;
border: 0;
font-size: 100%;
font: inherit;
vertical-align: baseline;
}

div { user agent stylesheet
display: block;

Inherited from div.content_list

a, abbr, common.css? v=2.411190841403:1
acronym, address, applet, article, aside,
audio, b, big, blockquote, body, canvas,
caption, center, cite, code, dd, del,
details, dfn, div, dl, dt, em, embed,
fieldset, figcaption, figure, footer, form,
h1, h2, h3, h4, h5, h6, header, hgroup, html,
i, iframe, img, ins, kbd, label, legend, li,
mark, menu, nav, object, ol, output, p, pre,
q, ruby, s, samp, section, small, span,
strike, strong, sub, summary, sup, table,
tbody, td, tfoot, th, thead, time, tr, tt, u,
ul, var, video {
margin: 0;
padding: 0;
border: 0;
font-size: 100%;
font: inherit;
vertical-align: baseline;
}

Inherited from div.content_article

a, abbr, common.css? v=2.411190841403:1
acronym, address, applet, article, aside,
audio, b, big, blockquote, body, canvas,

Web crawler: web pages

已为您找到 23968 套 上海租房

综合排序 最新上架 价格 面积

 a | 355.95 × 25
整租 · 武定西路1371弄 1室0厅 4500元
长宁-镇宁路 / 33m² / 南 / 1室0厅1卫
链家
① 1个月前发布
近地铁 随时看房


整租 · 花苑茶花园 2居室 5000
徐汇-康健 / 59m² / 南 / 2室1厅1卫
门店优选
① 5个月前发布
随时看房


整租 · 近8号线杨思站 简装两房 诚意出租
浦东-三林 / 61m² / 南 / 2室1厅1卫
链家
① 1个月前发布
近地铁 随时看房


整租 · 房东诚意出租 看房有钥匙 价格还能谈
徐汇-斜土路 / 57m² / 南北 / 2室1厅1卫
链家
① 1个月前发布

清空条件

Elements Console Sources Network Performance Memory Application Security Audits

Styles Computed Event Listeners

```
> div class="filter"> id="filter"
<div class="content w1150" id="content">
  <div class="content__article">
    <p class="content__title">
    <ul class="content__filter" id="contentList">
      <div class="content__list--item">
        <a class="content__list-item--aside" target="_blank" href="/zufang/SH2202640191251030016.html">
          
          <div class="content__list-item--main">
            <p class="content__list-item-title twoline">
              <a target="_blank" href="/zufang/SH2202640191251030016.html">整租 · 武定西路1371弄 1室0厅 4500元
            
            <div class="content__list-item--des">
              <a target="_blank" href="/zufang/changing">长宁
              "南"
              <a href="/zufang/zhenninglu" target="_blank">镇宁路
              <i>1室0厅1卫</i>
              <span class="hide"></span>
            
```

Web crawler: web pages

已为您找到 23968 套 上海租房

综合排序 最新上架 价格 面积

清空条件

整租 · 武定西路1371弄 1室0厅 4500元

长宁-镇宁路 / 33m² / 南 / 1室0厅1卫
链家
① 1个月前发布

 4500 元/月

近地铁 随时看房

整租 · 花苑茶花园 2居室 5000

徐汇-康健 / 59m² / 南 / 2室1厅1卫
门店优选
① 5个月前发布

 5000 元/月

随时看房

整租 · 近8号线杨思站 简装两房 诚意出租

浦东-三林 / 61m² / 南 / 2室1厅1卫
链家
① 1个月前发布

 4800 元/月

近地铁 随时看房

整租 · 房东诚意出租 看房有钥匙 价格还能谈

徐汇-斜土路 / 57m² / 南北 / 2室1厅1卫
链家
① 1个月前发布

 6500 元/月

Elements Console Sources Network Performance Memory Application Security Audits

Styles Computed Event Listeners

Filter element.style { }

```
body { index.css? v=20_411190841403:1 height: auto; position: relative; }
```

```
body { common.css? v=2_411190841403:1 font-size: 14px; color: #363636; font-family: HiraginoSansGB-W3,"Hiragino Sans GB","Microsoft Yahei UI","Microsoft Yahei","微软雅黑","Segoe UI",Tahoma,"宋体b80\4f53",SimSun,sans-serif; text-rendering: optimizeLegibility; -ms-text-size-adjust: 100%; -webkit-text-size-adjust: 100%; -webkit-font-smoothing: antialiased; -moz-osx-font-smoothing: grayscale; }
```

```
body, html { common.css? v=2_411190841403:1 background-color: #ffff; }
```

```
body { common.css? v=2_411190841403:1 line-height: 1; }
```

```
a, abbr, common.css? v=2_411190841403:1 acronym, address, applet, article, aside, audio, b, big, blockquote, body, canvas, caption, center, cite, code, dd, del, details, dfn, div, dl, dt, em, embed, fieldset, figcaption, figure, footer, form, h1, h2, h3, h4, h5, h6, header, hgroup, html, i, iframe, img, ins, kbd, label, legend, li, mark, menu, nav, object, ol, output, p, pre, q, ruby, s, samp, section, small, span, strike, strong, sub, summary, sup, table, tbody, td, tfoot, th, thead, time, tr, tt, u, ul, var, video { margin: 0; padding: 0; border: 0; font-size: 100%; font: inherit; vertical-align: baseline; }
```

```
body { user agent stylesheet display: block; margin-top: 8px; }
```

Inherited from `html`

```
a, abbr, common.css? v=2_411190841403:1 acronym, address, applet, article, aside, audio, b, big, blockquote, body, canvas, caption, center, cite, code, dd, del, details, dfn, div, dl, dt, em, embed,
```

Web crawler: web pages

- Download the web page.
 - import urllib.request
 - i=1
 - html =
`urllib.request.urlopen('http://sh.lianjia.com/zufang/d'+str(i)).read().decode('utf-8')`
- Try:
 - Go to <https://sh.lianjia.com/zufang/d1>
 - Download <https://sh.lianjia.com/zufang/d2>, …, up to <https://sh.lianjia.com/zufang/d100> by a for-loop



Web crawler: web pages

- Print the page and see what is in the html file.

```
In [4]: i=1  
html = urllib.request.urlopen('http://sh.lianjia.com/zufang/d'+str(i)).read().decode('utf-8')  
  
In [5]: print(html)
```

What is this?

```
<p class="content_list-item-title twoLine">  
    <a target="_blank" href="/zufang/SH2202640191251030016.html">  
        整租 · 武定西路1371弄 1室0厅 4500元  
    </a>  
</p>  
<p class="content_list-item-des">  
    <a target="_blank" href="/zufang/changning">长宁</a>-<a href="/zufang/zhenninglu" target="_  
    blank">镇宁路</a>  
    <i></i>  
    33m2  
    <i></i>南  
    1室0厅1卫  
    <i></i>  
    中楼层  
    <span class="hide">  
        (6层)  
    </span>  
</p>  
    <p class="content_list-item-brand oneline">  
        链家  
    </p>  
    <p class="content_list-item-time oneline">1个月前发布</p>  
<p class="content_list-item-bottom oneline">  
    <i class="content_item_tag-is_subway_house">近地铁</i>
```



Web crawler: web pages

- A detailed page is shown.

The screenshot shows a web browser displaying a real estate listing. The URL in the address bar is <https://sh.lianjia.com/zufang/SH2202640191251030016.html>. The page title is "整租 · 武定西路1371弄 1室0厅 4500元". Below the title, it says "房源上架时间 2019-03-07 房源编号: SH2202640191251030016" and there is a "举报" (Report) link. On the left, there is a floor plan diagram with dimensions: 2230 (width), 5105 (depth), and room areas: 卧室 15.6m², 过道 2.4m², 厨房 2.3m², and 卫生间 1.7m². To the right, a summary box displays the price "4500 元/月 (季付价)" in large red text, along with icons for "近地铁" (nearby subway), "随时看房" (viewing available), and other details: 整租 (whole lease), 1室0厅1卫 (1 bedroom, 0 living room, 1 bathroom), 33m² (33 square meters), and 朝南 (south-facing). A speaker icon in the bottom right corner indicates the page has audio content.

Web crawler: web pages

- Pieces of detailed information are listed in the URL with a string started with 'SH'.
- Next step is to obtain all the URLs such as <https://sh.lianjia.com/zufang/S...> and add them to a list.
- Knowledge of **regular expression** is needed.



Web crawler: regular expression

- A **regular expression**, **regex** or **regexp** is a sequence of characters that define a *search pattern*.
- Used to specify a set of strings required for a particular purpose. Often called a **pattern**.
- In 1950s mathematician Stephen Cole Kleene formalized the description of a *regular language*. It came into common use with Unix text-processing utilities in 80s.

Web crawler: regular expression

- **Boolean "or":**
 - A vertical bar separates alternatives. `gray|grey` can match "gray" or "grey".
- **Grouping:**
 - Parentheses are used to define the scope and precedence of the operators. `gray|grey` and `gr(a|e)y` are equivalent patterns, both describe the set of "gray" or "grey".

Web crawler: regular expression

- **Quantification:**
 - A quantifier after a token (e.g a character) or group specifies how often that a preceding element is allowed to occur. Common quantifiers are the question mark ?, the asterisk *, and the plus sign +.
- ?
 - Indicates zero or one occurrence of the preceding element. `colou?r` matches both "color" and "colour".
- *
 - Indicates zero or more occurrences of the preceding element. `ab*c` matches "ac", "abc", "abbc", "abbcc", and so on.



Web crawler: regular expression

- +
 - Indicates **one or more** occurrences of the preceding element. For example, ab+c matches "abc", "abbc", "abbbc", and so on, but not "ac".
- {n}
 - The preceding item is matched exactly n times.
- {min,}
 - The preceding item is matched min or more times.
- {min,max}
 - The preceding item is matched at least min times, but not more than max times.



Web crawler: regular expression

- **Wildcard**
 - The wildcard `.` matches any character. For example, `a.b` matches any string that contains an "a", then any other character and then a "b", `a.*b` matches any string that contains an "a" and a "b" at some later point.
 - These constructions can be combined to form **arbitrarily** complex expressions.

Web crawler: regular expression

- Package *re* in Python:
- *re.compile(pattern, flags=0)*
 - Compile a regular expression pattern into a regular expression object, to be used for matching using its `match()`, `search()` and other methods.
- *re.findall(pattern, string, flags=0)*
 - Return all non-overlapping matches of pattern in string, as a list of strings. The string is scanned left-to-right, and matches are returned in the order found. If one or more groups are present in the pattern, return a list of groups; this will be a list of tuples if the pattern has more than one group. Empty matches are included in the result.



Web crawler: regular expression

- Find URLs like /zufang/SH2202640191251030016.html in HTML

```
<div class="content_list-item-main">
  <p class="content_list-item-title twoline">
    <a target="_blank" href="/zufang/SH2202640191251030016.html">
      整租 · 武定西路1371弄 1室0厅 4500元
    </a>
  </p>
```



Web crawler: regular expression

```
import re
pattern_house_url = '<a target="_blank" href="(/zufang/SH\d*\.\html)">'
house_url = re.compile(pattern_house_url).findall(html)
```

- The highlight part is expected.

```
<div class="content_list-item-main">
    <p class="content_list-item-title twoline">
        <a target="_blank" href="/zufang/SH2202640191251030016.html">
            整租 · 武定西路1371弄 1室0厅 4500元
        </a>
    </p>
```



Web crawler: regular expression

- Print and check obtained URLs in the html.

```
In [7]: print(house_url)
```

```
['/zufang/SH2202640191251030016.html', '/zufang/SH2118065364868284416.html', '/zufang/SH2202735136326287360.htm  
l', '/zufang/SH2202827081300590592.html', '/zufang/SH2203600761156812800.html', '/zufang/SH2204255418761887744.h  
tml', '/zufang/SH2125351985594245120.html', '/zufang/SH2204278043474657280.html', '/zufang/SH220440016487240499  
2.html', '/zufang/SH2207071044735877120.html', '/zufang/SH2207104618712793088.html', '/zufang/SH2166375793389731  
840.html', '/zufang/SH2207731073776369664.html', '/zufang/SH2207837728912187392.html', '/zufang/SH22078874020798  
95552.html', '/zufang/SH2208692475068432384.html', '/zufang/SH2189736028540518400.html', '/zufang/SH220939232878  
1594624.html', '/zufang/SH2209407044606230528.html', '/zufang/SH2210088101043511296.html', '/zufang/SH2210185310  
271176704.html', '/zufang/SH2197570070966697984.html', '/zufang/SH2210826089419513856.html', '/zufang/SH22121732  
77810663424.html', '/zufang/SH2212203789283229696.html', '/zufang/SH2213061614511595520.html', '/zufang/SH213767  
0128563920896.html', '/zufang/SH2213082905108807680.html', '/zufang/SH2213685431558742016.html', '/zufang/SH2214  
222747201454080.html']
```

- These URLs are not complete.



Web crawler: regular expression

- Complete obtained URLs

- url_list = []
- for j in house_url:
- url_list.append('http://sh.lianjia.com'+str(j))

```
In [9]: print(url_list)
```

```
['http://sh.lianjia.com/zufang/SH2202640191251030016.html', 'http://sh.lianjia.com/zufang/SH2118065364868284416.html', 'http://sh.lianjia.com/zufang/SH2202735136326287360.html', 'http://sh.lianjia.com/zufang/SH2202827081300590592.html', 'http://sh.lianjia.com/zufang/SH2203600761156812800.html', 'http://sh.lianjia.com/zufang/SH2204255418761887744.html', 'http://sh.lianjia.com/zufang/SH2125351985594245120.html', 'http://sh.lianjia.com/zufang/SH2204278043474657280.html', 'http://sh.lianjia.com/zufang/SH2204400164872404992.html', 'http://sh.lianjia.com/zufang/SH2207071044735877120.html', 'http://sh.lianjia.com/zufang/SH2207104618712793088.html', 'http://sh.lianjia.com/zufang/SH2166375793389731840.html', 'http://sh.lianjia.com/zufang/SH2207731073776369664.html', 'http://sh.lianjia.com/zufang/SH2207837728912187392.html', 'http://sh.lianjia.com/zufang/SH2207887402079895552.html', 'http://sh.lianjia.com/zufang/SH2208692475068432384.html', 'http://sh.lianjia.com/zufang/SH2189736028540518400.html', 'http://sh.lianjia.com/zufang/SH2209392328781594624.html', 'http://sh.lianjia.com/zufang/SH2209407044606230528.html', 'http://sh.lianjia.com/zufang/SH2210088101043511296.html', 'http://sh.lianjia.com/zufang/SH2210185310271176704.html', 'http://sh.lianjia.com/zufang/SH2197570070966697984.html', 'http://sh.lianjia.com/zufang/SH2210826089419513856.html', 'http://sh.lianjia.com/zufang/SH2212173277810663424.html', 'http://sh.lianjia.com/zufang/SH2212203789283229696.html', 'http://sh.lianjia.com/zufang/SH2213061614511595520.html', 'http://sh.lianjia.com/zufang/SH2137670128563920896.html', 'http://sh.lianjia.com/zufang/SH2213082905108807680.html', 'http://sh.lianjia.com/zufang/SH2213685431558742016.html', 'http://sh.lianjia.com/zufang/SH2214222747201454080.html']
```



Web crawler: regular expression

- Visit the first URL in the url_list and get the html
- Use the pattern to get the content title from the html

```
In [10]: i=0
html = urllib.request.urlopen(url_list[i]).read().decode('utf-8')
```

```
In [11]: print(url_list[i])
http://sh.lianjia.com/zufang/S...1251030016.html
```


```
In [12]: pattern_introduction = '<p class="content_title">(.*)</p>'
introduction = re.compile(pattern_introduction)..findall(html)
print(introduction)
```

```
[‘整租 · 武定西路1371弄 1室0厅 4500元’]
```



Web crawler: regular expression

- More details to explore:
 - Room plan
 - Price
 - Period of payment
 - Agent
 -
- All can be found by **regular expression!**



Web crawler: beautiful soup

- [Beautiful Soup](#) is a Python library for pulling data out of HTML and XML files. It works with your favorite parser to provide idiomatic ways of navigating, searching, and modifying the parse tree. It commonly **saves programmers hours or days of work**.
- Let's try this new tool to grab something from
<https://sh.lianjia.com/zufang/S...>



Web crawler: beautiful soup

- Import packages.
 - from bs4 import BeautifulSoup
 - import requests
- To “fool” sh.lianjia.com, we need some mask.
 - header={'User-Agent':'Mozilla/5.0 (Macintosh; Intel Mac OS X 10_12_6) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/63.0.3239.108 Safari/537.36'}
- Download the page as usual.
 - data=requests.get('https://sh.lianjia.com/zufang/S^HH2202640191251030016.html',headers=header)
 - bs=BeautifulSoup(data.text)



Web crawler: beautiful soup

- We want to obtain the content title as we did previously.

```
In [32]: print(bs)

nloadQr?location=nav&amp;ljweb_channel_key=zufang_search"/>
</div>
</li>
</ul>
</div> <div class="search w1150" id="search">
<!-- <a class="search__logo" href="/" -->
<div class="search__wrap">
<input autocomplete="off" class="search__input fl" data-el="input" data-value="" placeholder="请输入区域、商圈或小区名开始找房" type="text" value="" />
<span class="search__button fl" data-el="button"></span>
</div>
</div> </div>
<!-- 房源有效时 -->
<div class="content clear w1150">
<!-- 房源标题 -->
<p class="content__title">整租 · 武定西路1371弄 1室0厅 4500元</p>
<!-- 房源副标题 -->
<div class="content__subtitle">
<i class="hide">4人浏览 </i>房源上架时间 2019-03-07
源编号: SH2202640191251030016</i>
<i class="house_code">房
```



Web crawler: beautiful soup

- Recall the pattern we wrote
 - pattern_introduction = '<p class="content_title">(.*)</p>'

```
In [32]: print(bs)

nloadQr?location=nav&amp;ljweb_channel_key=zufang_search"/>
</div>
</li>
</ul>
</div> <div class="search w1150" id="search">
<!-- <a class="search_logo" href="/" -->
<div class="search_wrap">
<input autocomplete="off" class="search_input fl" data-el="input" data-value="" placeholder="请输入区域、商圈或小区名开始找房" type="text" value="" />
<span class="search_button fl" data-el="button"></span>
</div>
</div> </div>
<!-- 房源有效时 -->
<div class="content clear w1150">
<!-- 房源标题 -->
<p class="content_title">整租 · 武定西路1371弄 1室0厅 4500元</p>
<!-- 房源副标题 -->
<div class="content_subtitle">
<i class="hide">4人浏览 </i>房源上架时间 2019-03-07
源编号: SH2202640191251030016</i>
<i class="house_code">房

```



Web crawler: beautiful soup

- We only need use *find* by specifying the tag `<p>` and class `'content_title'`
 - `bs.find('p', class_='content_title').text`

The screenshot shows a Jupyter Notebook interface. At the top, there is a snippet of HTML code. Below it, the 'In' cell contains a Python command to find the title using BeautifulSoup's `find` method. The 'Out' cell shows the resulting text '整租 · 武定西路1371弄 1室0厅 4500元'.

```
<p class="content_title">整租 · 武定西路1371弄 1室0厅 4500元</p>
<!-- 房源副标题 -->
<div class="content_subtitle">
<i class="hide">4人浏览 </i>房源上架时间 2019-03-07 <i class="house_code">房
源编号: SH202640191251030016</i>

In [33]: #<p class="content_title">整租 · 武定西路1371弄 1室0厅 4500元</p>
          bs.find('p', class_='content_title').text
Out[33]: '整租 · 武定西路1371弄 1室0厅 4500元'
```

We get the same result as the previously obtained.



Data mining starts from here!

- With obtained data, we start data mining.
- This part is organized as follows.
 - Data description
 - Visualization
 - Correlation



Data description

Variables	Description
url	url
title	Title
price	Price per month
housing_estate	Building
halls	Number of living rooms
rooms	Number of rooms
area	Total area
address	Address
price_per_area	Unit price (square meter)
district	District
block	Block
floor_type	Floor types (high, medium, low)
floor	Floor number

Data description

Variable	Description
<code>face_direction</code>	Room facing direction
<code>nearby_metro_line</code>	Metro line nearby
<code>nearby_metro_station</code>	Metro station nearby
<code>distance_to_metro</code>	Nearest Metro station
<code>rent_type</code>	Rental type (1 for 合租, 0 for 整租)
<code>has_own_washroom</code>	Private washroom availability (yes, no)
<code>has_own_balcony</code>	Private balcony availability (yes, no)
<code>is_on_sale</code>	On sale status (yes, no)
<code>update_time</code>	Update time
<code>bringin_records</code>	Bring-in records
<code>bringin_7_days</code>	Bring-in records within last 7 days
<code>bringin_count</code>	Number of bring-in records
<code>images</code>	Pictures of rooms for rent
<code>layout_image</code>	Floor plan
<code>time_to_pplsquare</code>	Hours travelling to people's square

Data description

- **pandas** is suited for different kinds of data:
 - **Tabular data** with heterogeneously-typed columns, as in an SQL table or Excel spreadsheet
 - Ordered and unordered (not necessarily fixed-frequency) **time series data**.
 - Arbitrary **matrix data** (homogeneously typed or heterogeneous) with row and column labels
 - Any other form of **observational / statistical data sets**. The data need not be labeled.

Data description

- Here are just a few of the things that pandas does well:
 - Handling missing data
 - Inserting and deleting columns from DataFrame
 - Aggregating and transforming data
 - Slicing, fancy indexing, and subsetting of large datasets
 - Reshaping of datasets
 -

Data description

Column/
variable



```
In [14]: import pandas as pd
```

```
In [19]: df_lj = pd.read_csv('house_data.csv')
```

```
In [21]: df_lj.head(3)
```

Out[21]:

	address	area	block	bringin_7_days	bringin_count	distance_to_metro	district	face_direction	fitment_type	floor	...	nearby_metro_station	price	pi	price	price_per_area	rent_type	rooms	title	update_time	url	time_to_ppisquare	Unnamed:
0	佳林路 115弄 (一 期) 佳 林路82 弄 (二 期) 佳 京路99 弄 (三 期) 金 高路 1617弄 (四 期)	171	金桥	0	2	0	浦东	(进门) 南	0	0	...	NaN	14500.0	14500.0	84.795322	0	4	链家好 房， 大气 四室	2017.01.13	http://sh.lianjia.com/zufang/shz3723680.html	-1	NaN	

Row/Observation



1	宋园路 69弄	110	古北	0	0	535	长宁	朝南北	0	16	...	宋园路	19000.0	19000.0	172.727273	0	2	圣美 邸， 有钥匙方 便 看， 南北 户型， 2室2 厅1 卫	2017.02.02	http://sh.lianjia.com/zufang/shz3739422.html	20	NaN
2	嘉涌路 99弄	43	江桥	2	2	0	嘉定	朝西	0	9	...	NaN	4000.0	4000.0	93.023256	0	2	MAX 未来， 好楼层， 放心好 房， 高清实拍	2017.02.02	http://sh.lianjia.com/zufang/shz3739516.html	-1	NaN

Data description

- To simplify the case, we select variables such as block, district, area, and price_per_area to study.
- Select desired columns.

```
→ In [36]: df_use = df_lj[['block', 'district', 'area', 'price_per_area', 'rent_type']]  
In [37]: df_use.head(5)  
Out[37]:  


|   | block | district | area | price_per_area | rent_type |
|---|-------|----------|------|----------------|-----------|
| 0 | 金桥    | 浦东       | 171  | 84.795322      | 0         |
| 1 | 古北    | 长宁       | 110  | 172.727273     | 0         |
| 2 | 江桥    | 嘉定       | 43   | 93.023256      | 0         |
| 3 | 万里    | 普陀       | 77   | 77.922078      | 0         |
| 4 | 曹路    | 浦东       | 89   | 65.168539      | 0         |


```

Data description

- First of all, we want to see some descriptive statistics of numerical variables.
- Use *df.describe* to obtain count, mean, standard deviation, and quantiles.

```
In [46]: df_use.describe()
Out[46]:
      area  price_per_area  rent_type
count  45182.000000  45182.000000  45182.000000
mean   83.676619    88.035169   0.074631
std    62.806153   174.941158   0.262799
min    1.000000    1.052632   0.000000
25%    50.000000   51.562500   0.000000
50%    76.000000   76.923077   0.000000
75%   102.000000  109.375000   0.000000
max   3800.000000  35000.000000  1.000000
```

Data description

- We now have 45182 candidates totally.
- If we want to know numbers of rental rooms **in each district**, we need to use *df.groupby*.
- We call this action **aggregation**.
- See details in
 - <https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.groupby.html>

```
In [42]: df_use.groupby(by='district').size()
```

```
Out[42]: district
上海周边      12
嘉定        2900
奉贤       1099
宝山       3348
崇明         2
徐汇       3155
普陀       2660
杨浦       2081
松江       2763
浦东     12283
虹口       1375
金山         20
长宁       2077
闵行       6271
闸北       1549
青浦       1342
静安         890
黄浦       1355
dtype: int64
```

Data description

- We obtain three more statistics for each district by applying the mean aggregation function for each numerical variables.
- What conclusion can we get for each of the means?
- Recall that to `rent_type`, we define 1 for 合租 and 0 for 整租.

```
In [67]: df_use.groupby(by='district').mean().reset_index()
```

Out[67]:

	district	area	price_per_area	rent_type
0	上海周边	111.333333	47.671525	0.166667
1	嘉定	77.673103	62.892753	0.100345
2	奉贤	87.767061	31.523732	0.006369
3	宝山	74.882616	68.835521	0.116189
4	崇明	103.500000	24.175610	0.000000
5	徐汇	75.036767	114.273836	0.043740
6	普陀	74.146617	94.756379	0.057143
7	杨浦	71.981259	99.239025	0.058626
8	松江	98.988780	64.142876	0.132103
9	浦东	86.412196	86.204175	0.081169
10	虹口	72.643636	106.209820	0.085091
11	金山	72.550000	161.646349	0.000000
12	长宁	84.701974	124.959281	0.028406
13	闵行	86.266784	75.251716	0.083400
14	闸北	68.352485	103.241463	0.078115
15	青浦	124.092399	77.410788	0.000000
16	静安	89.450562	152.795272	0.021348
17	黄浦	86.926937	143.029715	0.051661

Visualization


- Pyecharts: 30+ kinds of charts and maps of hundreds of regions supported.
- See the following link for more details.
 - <https://github.com/pyecharts/pyecharts>



Visualization

- scatter plot: area vs unit price.

- from pyecharts import Scatter
- v1 = df_use['area']
- v2 = df_use['price_per_area']
- scatter = Scatter("area vs unit price")
- scatter.add("", v1, v2)
- scatter.render()




Visualization

- Scatter plot is a usual way to plot raw data.
- What else can you plot?
- Visualize some aggregation results in various ways.

Visualization

- District vs mean area.
- df_mean =
df_use.groupby(by='district').mean().reset_index()
- from pyecharts import Bar
- bar = Bar('First plot','A bar plot of mean area')
- kwargs = dict(
 - name = 'Bar plot',
 - x_axis = df_mean['district'],
 - y_axis = df_mean['area'])
- bar.add(**kwargs)
- bar.render()
- With similar code, you can generate bar plots for the other two variables.



Visualization

- Word cloud is popular and size of words is controlled by area in this case.
- from pyecharts import WordCloud
- name_list = df_mean['district']
- value_list = df_mean['area']
- wordcloud = WordCloud(width=800, height=500)
- wordcloud.add("", name_list, value_list, word_size_range=[20, 100])
- wordcloud.render()



Visualization

- What is more?
- You may show data with Shanghai map involved.
- Go to <https://pyecharts.org/#/zh-cn/intro> and give a try.

Correlation

- Measure relationship between area and unit price
- **Correlation** refers to the degree of relationship (or dependency) between two variables.
- Linear correlation refers to straight-line relationships between two variables.
- A correlation can range between -1 (perfect negative relationship) and +1 (perfect positive relationship), with 0 indicating no straight-line relationship.

Pearson correlation


$$r_{xy} = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

where:

- n is sample size
- x_i, y_i are the individual sample points indexed with i
- $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i$ (the sample mean); and analogously for \bar{y}

Correlation

- Sub-figure (a) and (d) show perfect linear correlation.
- Sub-figure (f) shows zero correlation even though the pattern is obvious.



Correlation

- **NumPy**
 - N-dimensional array object
 - sophisticated (broadcasting) functions
 - tools for integrating C/C++ and Fortran code
 - useful linear algebra, Fourier transform, and random number capabilities

Correlation

- With **NumPy** function *corrcoef*, we can obtain the correlation between area and price_per_area.

- import numpy as np
- x = df_use['area'].values
- y = df_use['price_per_area'].values
- np.corrcoef(x,y)

- The result is called

Correlation matrix.

```
In [87]: import numpy as np
```

```
In [99]: x = df_use['area'].values  
y = df_use['price_per_area'].values  
np.corrcoef(x,y)
```

```
Out[99]: array([[ 1.        , -0.16117171],  
                 [-0.16117171,  1.        ]])
```

Linear regression

- In statistics, **linear regression** is a linear approach to modelling the *relationship* between a scalar response (or dependent variable) and one or more explanatory variables (or independent variables).
- In linear regression, the relationships are modeled using linear predictor functions whose unknown model parameters are estimated from the data.
- used extensively in practical applications.

Linear regression

- **scikit-learn** - *Machine Learning in Python*
 - Simple and efficient tools for data mining and data analysis
 - Accessible to everybody, and reusable in various contexts
 - Built on NumPy, SciPy, and matplotlib
 - Open source, commercially usable - BSD license

Linear regression



The screenshot shows the official scikit-learn website. At the top, there's a navigation bar with links for Home, Installation, Documentation, Examples, Google Custom Search, and a 'Fork me on GitHub' button. Below the navigation is a grid of 27 small plots demonstrating various machine learning models like Logistic Regression, Linear SVM, and Ridge SVM.

Classification
Identifying to which category an object belongs to.
Applications: Spam detection, Image recognition.
Algorithms: SVM, nearest neighbors, random forest, ...
[— Examples](#)

Regression
Predicting a continuous-valued attribute associated with an object.
Applications: Drug response, Stock prices.
Algorithms: SVR, ridge regression, Lasso, ...
[— Examples](#)

Clustering
Automatic grouping of similar objects into sets.
Applications: Customer segmentation, Grouping experiment outcomes
Algorithms: k-Means, spectral clustering, mean-shift, ...
[— Examples](#)

Dimensionality reduction
Reducing the number of random variables to consider.
Applications: Visualization, Increased efficiency
Algorithms: PCA, feature selection, non-negative matrix factorization.
[— Examples](#)

Model selection
Comparing, validating and choosing parameters and models.
Goal: Improved accuracy via parameter tuning
Modules: grid search, cross validation, metrics.
[— Examples](#)

Preprocessing
Feature extraction and normalization.
Application: Transforming input data such as text for use with machine learning algorithms.
Modules: preprocessing, feature extraction.
[— Examples](#)

Linear regression

- analyze the relationship between area and price_per_area using *linear_model* in **sklearn**.
 - import matplotlib.pyplot as plt
 - import numpy as np
 - from sklearn import linear_model
 - regr = linear_model.LinearRegression()
 - x = df_use['area'].values.reshape(-1, 1)
 - y = df_use['price_per_area'].values
 - regr.fit(x, y)

Linear regression

- The linear model price_per_area vs area is built.
- Price_per_area = 100.19-0.15*area

```
In [119]: print('Coefficients: \n', regr.coef_, '\n Intercept: \n', regr.intercept_)
```

```
Coefficients:  
[-0.15454709]  
Intercept:  
100.19471649716381
```

- This equation tells us that price_per_area gets 0.15 Yuan cheaper as area gets 1 square meter larger.

Linear regression

- We can plot the model.
- We should mention that analysis here is not rigorous at all. The statistical knowledge needed is out of the scope of the course.

```
In [120]: y_pred = regr.predict(x)
```

```
In [121]: # Plot outputs  
plt.scatter(x, y, color='black')  
plt.plot(x, y_pred, color='blue', linewidth=3)  
  
plt.xticks([])  
plt.yticks([])  
  
plt.show()
```



Thanks!